

Homework4

November 24, 2022

1 Question 1

1.0.1 For the prostate data of Chapter 3 (ESL), carry out a best-subset linear regression analysis, as in Table 3.3 (third column from the left). Compute the AIC, BIC, five- and tenfold cross-validation, and bootstrap .632 estimates of prediction error. The data can be obtained from the book website.

```
[357]: install.packages("ISLR")
install.packages("knitr")
install.packages("printr")
install.packages('leaps')
install.packages('bootstrap')
install.packages('lattice')
install.packages('caret')
library(lattice)
library(caret)
library(bootstrap)
library(leaps)
library(ISLR)
library(knitr)
library(printr)
library(boot)
```

The downloaded binary packages are in
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages
also installing the dependency 'xfun'

The downloaded binary packages are in
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

The downloaded binary packages are in
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

The downloaded binary packages are in

```
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages
```

The downloaded binary packages are in

```
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages
```

The downloaded binary packages are in

```
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages
```

The downloaded binary packages are in

```
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages
```

Warning message in FUN(X[[i]], ...):

```
"unknown type in R_decompress3"
```

```
Error in FUN(X[[i]], ...): lazy-load database '/Library/Frameworks/R.framework/  
↳ Versions/4.2/Resources/library/lattice/data/Rdata.rdb' is corrupt
```

```
Traceback:
```

```
1. library(boot)  
2. checkConflicts(package, pkgname, pkgpath, nogenerics, ns)  
3. same.isFn(i)  
4. vapply(same, exists, NA, where = where, mode = "function", inherits = FALSE)  
5. FUN(X[[i]], ...)
```

load the data

```
[21]: install.packages("ElemStatLearn_2015.6.26.tar.gz", repos = NULL, type =  
↳ "source")  
library("ElemStatLearn")  
data(prostate)
```

Warning message in install.packages("ElemStatLearn_2015.6.26.tar.gz", repos =
NULL, :

```
"installation of package 'ElemStatLearn_2015.6.26.tar.gz' had non-zero exit  
status"
```

```
[22]: str( prostate )
```

```
'data.frame': 97 obs. of 10 variables:  
 $ lcavol : num -0.58 -0.994 -0.511 -1.204 0.751 ...  
 $ lweight: num 2.77 3.32 2.69 3.28 3.43 ...  
 $ age : int 50 58 74 58 62 50 64 58 47 63 ...  
 $ lbph : num -1.39 -1.39 -1.39 -1.39 -1.39 ...  
 $ svi : int 0 0 0 0 0 0 0 0 0 0 ...  
 $ lcp : num -1.39 -1.39 -1.39 -1.39 -1.39 ...  
 $ gleason: int 6 6 7 6 6 6 6 6 6 6 ...  
 $ pgg45 : int 0 0 20 0 0 0 0 0 0 0 ...  
 $ lpsa : num -0.431 -0.163 -0.163 -0.163 0.372 ...
```

```
$ train : logi TRUE TRUE TRUE TRUE TRUE TRUE ...
```

Computing best subsets regression

```
[29]: train_set<-subset(prostate,train=="TRUE")[,1:9]
test_set<-subset(prostate,train=="FALSE")[,1:9]
y_train<-train_set$lpsa
y_test<-test_set$lpsa

regfit.full<-regsubsets(lpsa~.,data=train_set,
  ↪nbest=1,nvmax=8,method="exhaustive")
my_sum <- summary(regfit.full)
summary(regfit.full)
```

Subset selection object

Call: regsubsets.formula(lpsa ~ ., data = train_set, nbest = 1, nvmax = 8,
method = "exhaustive")

8 Variables (and intercept)

Forced in Forced out

lcavol	FALSE	FALSE
lweight	FALSE	FALSE
age	FALSE	FALSE
lbph	FALSE	FALSE
svi	FALSE	FALSE
lcp	FALSE	FALSE
gleason	FALSE	FALSE
pgg45	FALSE	FALSE

1 subsets of each size up to 8

Selection Algorithm: exhaustive

		lcavol	lweight	age	lbph	svi	lcp	gleason	pgg45
1	(1)	"*"	" "	" "	" "	" "	" "	" "	" "
2	(1)	"*"	"*"	" "	" "	" "	" "	" "	" "
3	(1)	"*"	"*"	" "	" "	"*"	" "	" "	" "
4	(1)	"*"	"*"	" "	"*"	"*"	" "	" "	" "
5	(1)	"*"	"*"	" "	"*"	"*"	" "	" "	"*"
6	(1)	"*"	"*"	" "	"*"	"*"	"*"	" "	"*"
7	(1)	"*"	"*"	"*"	"*"	"*"	"*"	" "	"*"
8	(1)	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"

1.1 Model selection criteria: AIC & BIC

```
[53]: select = my_sum$outmat
train_error_store<-c()
test_error_store<-c()
AIC_store<-c()
BIC_store<-c()

for(i in 1:8){
```

```

temp<-which(select[i,]=="*")
red_train<-train_set[,c(9,temp)]
red_test<-test_set[,c(9,temp)]
red_fit<-lm(lpsa~.,data=red_train)
AIC<-AIC(red_fit)
BIC<-BIC(red_fit)
predict_train<-predict(red_fit,newdata=red_train)
predict_test<-predict(red_fit,newdata=red_test)
train_error<-sum((predict_train-y_train)^2)/length(y_train)
test_error<-sum((predict_test-y_test)^2)/length(y_test)
train_error_store<-c(train_error_store,train_error)
test_error_store<-c(test_error_store,test_error)
AIC_store<-c(AIC_store,AIC)
BIC_store<-c(BIC_store,BIC)

}

# print('Train error')
# train_error_store
# min(train_error_store)
# print('Test error')
# test_error_store
# min(test_error_store)

print(paste("Subset selected by AIC =",which.min(AIC_store)))
print('AIC values')
AIC_store

print(paste("Subset selected by BIC =",which.min(BIC_store)))
print('BIC values')
BIC_store

print(paste('Test error of 2 variable model:',test_error_store[2]))

upper= max(AIC_store,BIC_store)
lower= min(AIC_store,BIC_store)

plot(AIC_store,type="o",lty=2,col = "blue",ylim = c(lower-5,upper+5),xlab = "k",
     ↪ "k",main="AIC & BIC",ylab="Value")
lines(BIC_store,type="o",lty=1,col="red")
legend("topright",c("AIC", "BIC"),lty=c(2,1),col=c("blue","red"))

```

```

plot(train_error_store,type="o",lty=2,col = "blue",ylim = c(0,1),xlab = "k",ylab="error",main="Error")
lines(test_error_store,type="o",lty=1,col="red")
legend("topright",c("training.error", "test.error"),lty=c(2,1),col=c("blue","red"))

```

[1] "Subset selected by AIC = 7"

[1] "AIC values"

1. 168.764154371389 2. 158.520967079593 3. 156.454849884341 4. 154.31269099965
5. 154.772911236327 6. 153.498370284356 7. 153.03495137847 8. 155.010102017978


[1] "Subset selected by BIC = 2"

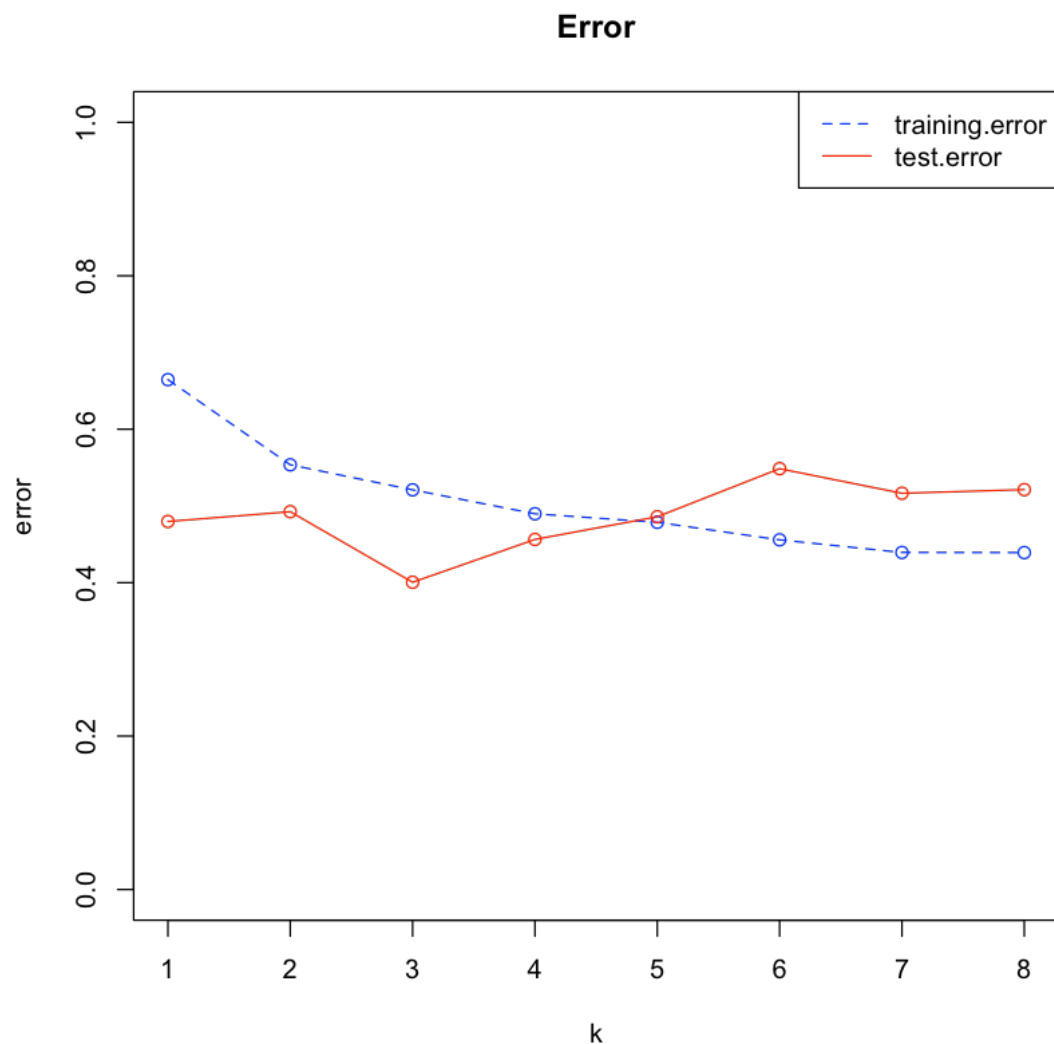
[1] "BIC values"

1. 175.378232229561 2. 167.339737557157 3. 167.478312981296 4. 167.540846715996
5. 170.205759572064 6. 171.135911239484 7. 172.877184952989 8. 177.057028211888

[1] "Test error of 2 variable model: 0.492482349035638"

AIC & BIC





The model with the lowest AIC is the best among several. We discover a broader model with seven factors for the AIC (lcavol, lweight, age, lbph, svi, lcp ,pgg45). Only the predictor Gleason was eliminated during this selection process. Since the BIC penalizes the number of parameters more severely, it chooses smaller models. In this instance, we discover that the ideal model includes lweight and lcavol.

We can observe that this model's test error is 0.49.

1.2 Model selection criteria: k-cross validation

1.2.1 k=5

```
[128]: predict.regsubsets = function(object,newdata,id,...){  
  form = as.formula(object$call[[2]]) # Extract the formula used when we  
  ↪called regsubsets()  
  mat = model.matrix(form,newdata)    # Build the model matrix  
  coefi = coef(object,id=id)          # Extract the coefficients of the ith  
  ↪model  
  xvars = names(coefi)                # Pull out the names of the  
  ↪predictors used in the ith model  
  mat[,xvars]%*%coefi                 # Make predictions using matrix  
  ↪multiplication  
}
```

```
[129]: k = 5          # number of folds  
set.seed(1)          # set the random seed so we all get the same results  
  
# Assign each observation to a single fold  
folds = sample(1:k, nrow(prostate), replace = TRUE)  
  
# Create a matrix to store the results of our upcoming calculations  
cv_errors = matrix(NA, k, 8, dimnames = list(NULL, paste(1:8)))
```

```
[130]: # Outer loop iterates over all folds  
for(j in 1:k){  
  
  # The perform best subset selection on the full dataset, minus the jth fold  
  best_fit = regsubsets(lpsa~., data = prostate[folds!=j,], nvmax=8)  
  
  # Inner loop iterates over each size i  
  for(i in 1:8){  
  
    # Predict the values of the current fold from the "best subset" model  
    ↪on i predictors  
    pred = predict(best_fit, prostate[folds==j,], id=i)  
  
    # Calculate the MSE, store it in the matrix we created above  
    cv_errors[j,i] = mean((prostate$lpsa[folds==j]-pred)^2)  
  }  
}
```


```
[149]: # Take the mean of over all folds for each model size  
k5_mean_cv_errors = apply(cv_errors, 2, mean)  
  
# Find the model size with the smallest cross-validation error  
min = which.min(k5_mean_cv_errors)
```



```
# Plot the cross-validation error for each model size, highlight the min
plot(k5_mean_cv_errors, type='b')
points(min, k5_mean_cv_errors[min][1], col = "red", cex = 2, pch = 20)

print(paste("MSE for k=5: ", min(mean_cv_errors)))
```

```
[1] "MSE for k=5: 0.51668234608548"
```



```
[150]: reg_best = regsubsets(lpsa~., data = prostate, nvmax = 8)
coef(reg_best, 3)
```

```
(Intercept) -0.777156641580076 lcavol 0.52585188198094 lweight 0.661769911594472 svi
```

0.665666562857202

1.2.2 k=10

```
[151]: k = 10          # number of folds
set.seed(567)      # set the random seed so we all get the same results

# Assign each observation to a single fold
folds = sample(1:k, nrow(prostate), replace = TRUE)

# Create a matrix to store the results of our upcoming calculations
cv_errors = matrix(NA, k, 8, dimnames = list(NULL, paste(1:8)))

[152]: # Outer loop iterates over all folds
for(j in 1:k){

  # The perform best subset selection on the full dataset, minus the jth fold
  best_fit = regsubsets(lpsa~., data = prostate[folds!=j,], nvmax=8)

  # Inner loop iterates over each size i
  for(i in 1:8){

    # Predict the values of the current fold from the "best subset" model
    ↪ on i predictors
    pred = predict(best_fit, prostate[folds==j,], id=i)

    # Calculate the MSE, store it in the matrix we created above
    cv_errors[j,i] = mean((prostate$lpsa[folds==j]-pred)^2)
  }
}


[153]: # Take the mean of over all folds for each model size
k10_mean_cv_errors = apply(cv_errors, 2, mean)

# Find the model size with the smallest cross-validation error
min = which.min(k10_mean_cv_errors)

# Plot the cross-validation error for each model size, highlight the min
plot(k10_mean_cv_errors, type='b')
points(min, k10_mean_cv_errors[min][1], col = "red", cex = 2, pch = 20)

print(paste("MSE for k=10: ",min(mean_cv_errors)))
```

```
[1] "MSE for k=10: 0.51668234608548"
```



```
[154]: reg_best = regsubsets(lpsa~., data = prostate, nvmax = 8)
coef(reg_best, 3)
```

```
(Intercept)  -0.777156641580076 lcavol   0.52585188198094 lweight   0.661769911594472 svi
0.665666562857202
```

Above, we implemented best-subset cross-validation for linear regression. We call that procedure and then receive the set of predictors from the $k=5$, $k=10$ that have the smallest cross-validated estimate of the mean square error. Similar subset with three variables is obtained using 5-fold and 10-fold cross-validation (lcavol,lweight,svi).

1.3 Model selection criteria: bootstrap.632

```
[156]: set.seed(1)

x<-prostate[,1:8]
y<-prostate[,9]

theta_fit<-function(x,y){lsfit(x,y)}
theta_predict<-function(fit,x){cbind(1,x)%*%fit$coef}
sq_err<-function(y,yhat){(y-yhat)^2}

bootstrap_632_error_store<-c()
for(i in 1:8){
  temp<-which(select[i,]=="*")
  res<-bootpred(x[,temp],y,nboot = 50,theta_fit,theta_predict,err.meas=sq_err)
  bootstrap_632_error_store<-c(bootstrap_632_error_store,res[[3]])
}


bootstrap_632_error_store
print(paste("Bootstrp.632 error for ",which.
  ↪min(bootstrap_632_error_store),"variabl model_
  ↪is",min(bootstrap_632_error_store)))

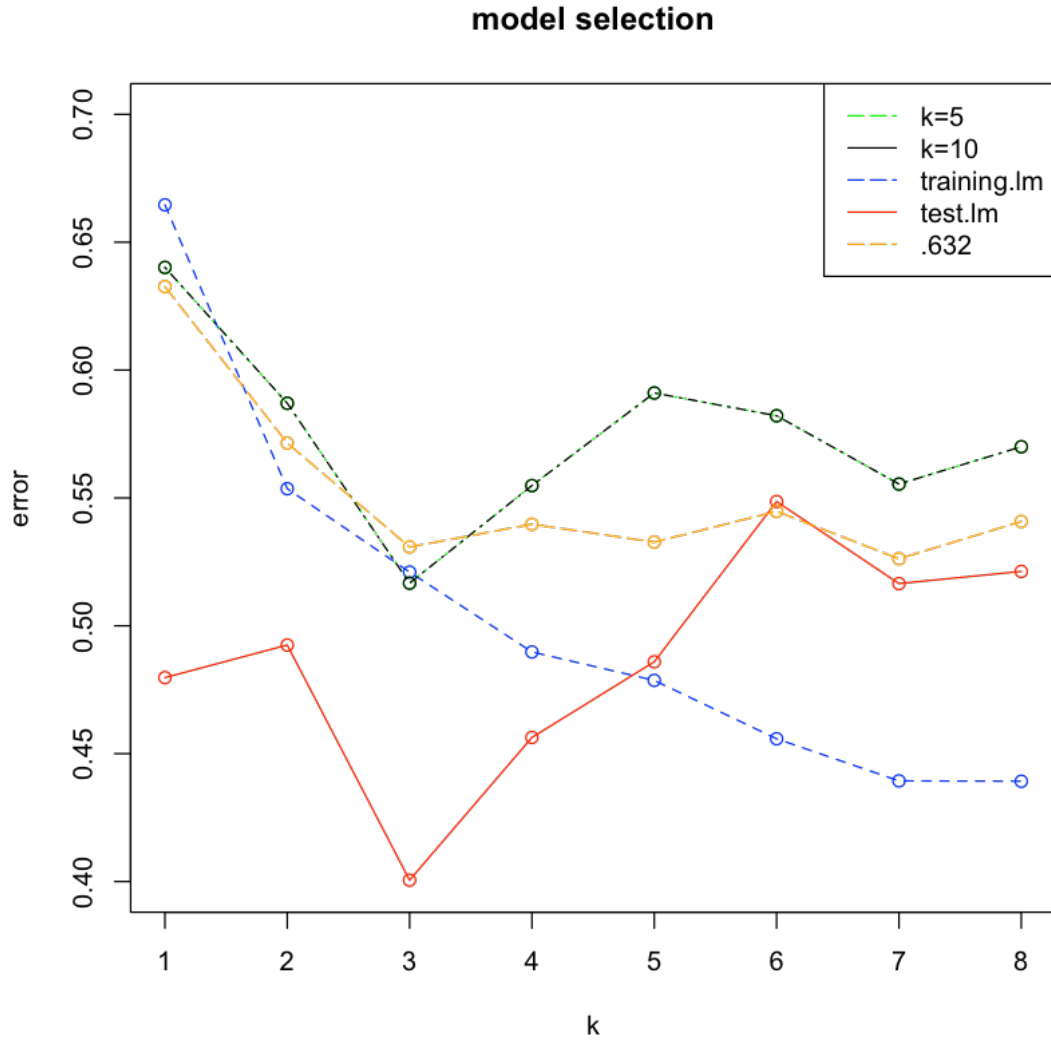
plot(bootstrap_632_error_store,type="o",lty=5,col="orange",main="error",xlab =_
  ↪"k",ylab="error")
legend("topright",c(".632"),lty=1,col=c("orange"))

plot(train_error_store,type="o",lty=2,col = "blue",ylim = c(0.4,0.7),xlab =_
  ↪"k",ylab="error",main="model selection")
lines(test_error_store,type="o",lty=1,col="red")
lines(k5_mean_cv_errors,type="o",lty=3,col = "green")
lines(k10_mean_cv_errors,type="o",lty=4,col="black")
lines(bootstrap.632.error.store,type="o",lty=5,col="orange")
legend("topright",c("k=5", "k=10","training.lm", "test.lm",".
  ↪632"),lty=c(5,1),col=c("green","black","blue","red","orange"))
```

```
1. 0.632630056057999 2. 0.57146399530187 3. 0.530823037892348 4. 0.53971847264767
5. 0.53273057426127 6. 0.544743393174272 7. 0.526214505661281 8. 0.540796707029031
```

```
[1] "Bootstrp.632 error for 7 variabl model is 0.526214505661281"
```





The model with two variables is the best fit for this data set, according to the subset selection approach, which was based on BIC. The model with three variables appears to have the best accuracy, according to the MSE obtained by running Cross Validation (both $k=5$ and $k=10$) on the same data set. In contrast, the bootstrap.632 resampling approach estimates an error of 0.52 for a model with 7 variables that is similar to the AIC.

The 0.632 estimator uses bootstrap samples to calculate the prediction error. It gives equal weight to a bootstrapped estimate and the training error. The best subset size determined by different methods is 3, and the stated error lies between 0.50 and 0.51. The testing MSE of 0.49 appears to be a bit optimistic based on these figures.

2 Question 2

2.0.1 The Bikeshare data (ISLR2) contains the hourly and daily count of rental bikes between 2011-2012 in a bikeshare program. Other important features such as weather, and seasonal information, are also included. You are asked to construct a regression tree to predict the daily count of rental bikes.

2.0.2 (a) *Divide the data into test and training and perform model selection to determine the optimal tree size. Comment on the performance of the tree. How many test samples are assigned to each terminal region?*

```
[24]: install.packages("rpart")
install.packages('MASS')
install.packages('ISLR2')
install.packages('rpart.plot')
library(rpart.plot)
library(ISLR2)
library(rpart)
library(MASS)
```

The downloaded binary packages are in

/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

The downloaded binary packages are in

/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

There is a binary version available but the source version is later:

```
      binary source needs_compilation
ISLR2  1.3-1  1.3-2                  FALSE
```

installing the source package 'ISLR2'

The downloaded binary packages are in

/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

```
[229]: set.seed(2255)

bikeshare_df<-read.csv("/Users/sreeragvenugopalan/Desktop/Sem 1/Statistical_
↳Data Mining/HW/4/Bike-Sharing-Dataset/day.csv")
head(bikeshare_df,5)
```

		instant <int>	dteday <chr>	season <int>	yr <int>	mnth <int>	holiday <int>	weekday <int>	workingday <int>	weather <int>
A data.frame: 5 × 16	1	1	2011-01-01	1	0	1	0	6	0	2
	2	2	2011-01-02	1	0	1	0	0	0	2
	3	3	2011-01-03	1	0	1	0	1	1	1
	4	4	2011-01-04	1	0	1	0	2	1	1
	5	5	2011-01-05	1	0	1	0	3	1	1

2.0.3 PreProcess Data

We can see that column 'cnt' is the sum of 'casual' and 'registered', So we will remove those two columns.

```
[230]: bikeshare_df<-subset(bikeshare_df,select=-c(casual,registered))
       head(bikeshare_df,5)
```

		instant <int>	dteday <chr>	season <int>	yr <int>	mnth <int>	holiday <int>	weekday <int>	workingday <int>	weather <int>
A data.frame: 5 × 14	1	1	2011-01-01	1	0	1	0	6	0	2
	2	2	2011-01-02	1	0	1	0	0	0	2
	3	3	2011-01-03	1	0	1	0	1	1	1
	4	4	2011-01-04	1	0	1	0	2	1	1
	5	5	2011-01-05	1	0	1	0	3	1	1

Renaming the column titles.

```
[231]: names(bikeshare_df)<-c('ID','datetime','season','year','month','holiday','weekday','workingday','weather')
       head(bikeshare_df,5)
```

		ID <int>	datetime <chr>	season <int>	year <int>	month <int>	holiday <int>	weekday <int>	workingday <int>	weather <int>
A data.frame: 5 × 14	1	1	2011-01-01	1	0	1	0	6	0	2
	2	2	2011-01-02	1	0	1	0	0	0	2
	3	3	2011-01-03	1	0	1	0	1	1	1
	4	4	2011-01-04	1	0	1	0	2	1	1
	5	5	2011-01-05	1	0	1	0	3	1	1

Typcast some columns

```
[232]: bikeshare_df$datetime<- as.Date(bikeshare_df$datetime)
       bikeshare_df$year<-as.factor(bikeshare_df$year)
       bikeshare_df$month<-as.factor(bikeshare_df$month)
       bikeshare_df$season <- as.factor(bikeshare_df$season)
       bikeshare_df$holiday<- as.factor(bikeshare_df$holiday)
       bikeshare_df$weekday<- as.factor(bikeshare_df$weekday)
       bikeshare_df$workingday<- as.factor(bikeshare_df$workingday)
       bikeshare_df$weather_condition<- as.factor(bikeshare_df$weather_condition)
       head(bikeshare_df,5)
```


		ID <int>	datetime <date>	season <fct>	year <fct>	month <fct>	holiday <fct>	weekday <fct>	workingday <fct>	weather <fct>
A data.frame: 5 × 14	1	1	2011-01-01	1	0	1	0	6	0	2
	2	2	2011-01-02	1	0	1	0	0	0	2
	3	3	2011-01-03	1	0	1	0	1	1	1
	4	4	2011-01-04	1	0	1	0	2	1	1
	5	5	2011-01-05	1	0	1	0	3	1	1

check for NA values

```
[233]: any(is.na(bikeshare_df))
```

FALSE

Split the dataset into train and test

```
[234]: train_index<-sample(1:nrow(bikeshare_df),0.7*nrow(bikeshare_df))
train_data<-bikeshare_df[train_index,]
test_data<-bikeshare_df[-train_index,]
head(train_data,5)
head(test_data,5)
```

		ID <int>	datetime <date>	season <fct>	year <fct>	month <fct>	holiday <fct>	weekday <fct>	workingday <fct>	weather <fct>
A data.frame: 5 × 14	29	29	2011-01-29	1	0	1	0	6	0	1
	557	557	2012-07-10	3	1	7	0	2	1	2
	689	689	2012-11-19	4	1	11	0	1	1	2
	117	117	2011-04-27	2	0	4	0	3	1	2
	541	541	2012-06-24	3	1	6	0	0	0	1

		ID <int>	datetime <date>	season <fct>	year <fct>	month <fct>	holiday <fct>	weekday <fct>	workingday <fct>	weather <fct>
A data.frame: 5 × 14	2	2	2011-01-02	1	0	1	0	0	0	2
	7	7	2011-01-07	1	0	1	0	5	1	2
	13	13	2011-01-13	1	0	1	0	4	1	1
	20	20	2011-01-20	1	0	1	0	4	1	2
	23	23	2011-01-23	1	0	1	0	0	0	1

```
[235]: train<-subset(train_data,select=c('season','year','month','holiday',
↪'weekday','workingday','weather_condition','temp','humidity','windspeed','total_count'))
test<-subset(test_data,select=c('season','year','month','holiday','weekday','workingday','weather_condition'))
```

One-hot encoding the categorical values using dummyVars

```
[236]: dumy_cat_vals <- dummyVars('~
↪season+holiday+workingday+weather_condition+year', data = train)
dumy_cat_vals_df <- data.frame(predict(dumy_cat_vals, newdata = train))
```

create a new subset for train and test; categorical, numerical attributes

```
[237]: train_cat_attributes<-subset(train,select=c('season','holiday','workingday','weather_condition',
test_cat_attributes<-subset(test,select=c('season','holiday','workingday','weather_condition',
train_num_attributes<-subset(train,select=c('weekday','month','temp','humidity','windspeed','total_count',
test_num_attributes<-subset(test,select=c('weekday','month','temp','humidity','windspeed','total_count'))
```

```
[238]: train_encoded_attributess<-cbind(train_num_attributes,dummy_cat_vals_df)
head(train_encoded_attributess,5)
```

		weekday	month	temp	humidity	windspeed	total_count	season.1	season.2
		<fct>	<fct>	<dbl>	<dbl>	<dbl>	<int>	<dbl>	<dbl>
	29	6	1	0.196522	0.651739	0.145365	1098	1	0
A data.frame: 5 × 19	557	2	7	0.720833	0.667500	0.151737	6290	0	0
	689	1	11	0.380833	0.623333	0.235067	5499	0	0
	117	3	4	0.620000	0.835417	0.312200	3872	0	1
	541	0	6	0.743333	0.479167	0.145525	6891	0	0

Building and training the model

```
[239]: rpart.control<-rpart.control(minbucket = 2,minsplit = 4, xval = 10, cp = 0.01)
original_fit_bshare<-rpart(train_encoded_attributess$total_count~.
, data=train_encoded_attributess[, -c(6)], control=rpart.
control, method='anova', cp=0.01)
original_fit_bshare
```

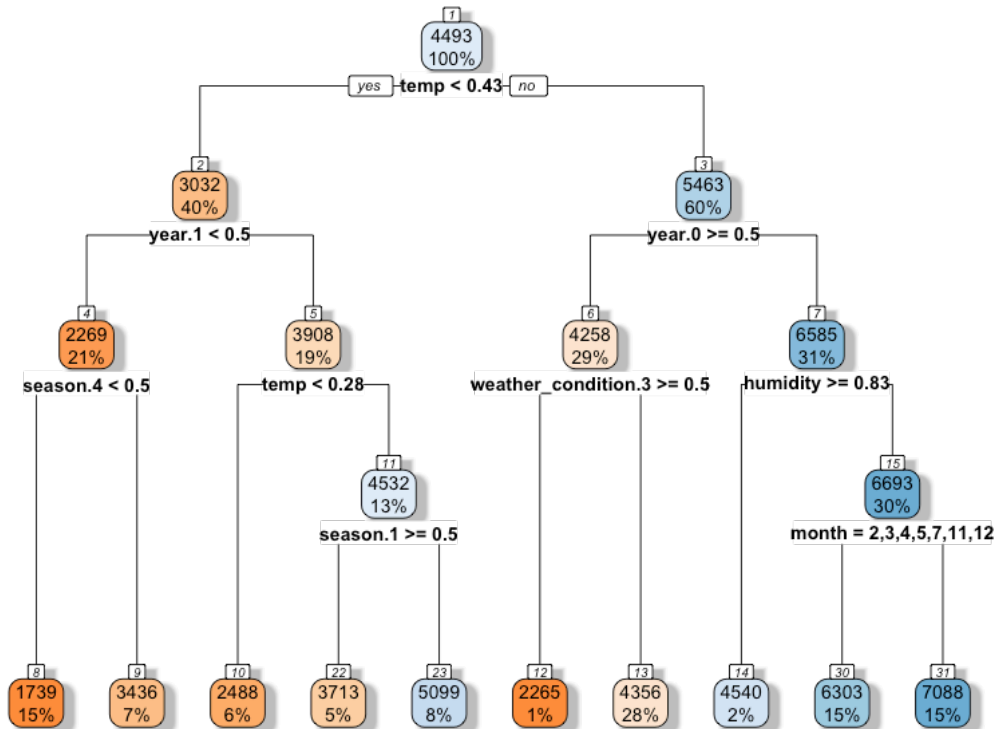
n= 511

node), split, n, deviance, yval
* denotes terminal node

```
1) root 511 1882115000.0 4492.544
  2) temp< 0.432373 204 449964900.0 3032.240
    4) year.1< 0.5 109 114334800.0 2268.734
      8) season.4< 0.5 75 29791150.0 1739.467 *
      9) season.4>=0.5 34 17190340.0 3436.235 *
    5) year.1>=0.5 95 199185000.0 3908.263
      10) temp< 0.2804165 29 14609140.0 2488.448 *
      11) temp>=0.2804165 66 100428300.0 4532.121
        22) season.1>=0.5 27 20318370.0 3713.000 *
        23) season.1< 0.5 39 49452270.0 5099.205 *
  3) temp>=0.432373 307 708048800.0 5462.909
    6) year.0>=0.5 148 109593600.0 4257.520
      12) weather_condition.3>=0.5 7 858316.9 2264.857 *
      13) weather_condition.3< 0.5 141 79560490.0 4356.447 *
    7) year.0< 0.5 159 183255500.0 6584.906
      14) humidity>=0.8322915 8 2719668.0 4539.750 *
      15) humidity< 0.8322915 151 145301700.0 6693.258
        30) month=2,3,4,5,7,11,12 76 65508410.0 6303.447 *
```

31) month=6,8,9,10 75 56542530.0 7088.267 *

```
[240]: rpart.plot(original_fit_bshare, box.palette="OrBu", shadow.col="gray",
  ↪nn=TRUE,roundint=FALSE)
```



```
[275]: min_cp = which.min(original_fit_bshare$cptable[,4])
  pruned_fit_bshare <- prune(original_fit_bshare, cp =
  ↪original_fit_bshare$cptable[min_cp, 1])
  pruned_fit_bshare
```

n= 511

node), split, n, deviance, yval


* denotes terminal node

```
1) root 511 1882115000.0 4492.544
 2) temp< 0.432373 204 449964900.0 3032.240
    4) year.1< 0.5 109 114334800.0 2268.734
      8) season.4< 0.5 75 29791150.0 1739.467 *
      9) season.4>=0.5 34 17190340.0 3436.235 *
    5) year.1>=0.5 95 199185000.0 3908.263
      10) temp< 0.2804165 29 14609140.0 2488.448 *
      11) temp>=0.2804165 66 100428300.0 4532.121
        22) season.1>=0.5 27 20318370.0 3713.000 *
        23) season.1< 0.5 39 49452270.0 5099.205 *
 3) temp>=0.432373 307 708048800.0 5462.909
    6) year.0>=0.5 148 109593600.0 4257.520
      12) weather_condition.3>=0.5 7 858316.9 2264.857 *
      13) weather_condition.3< 0.5 141 79560490.0 4356.447 *
 7) year.0< 0.5 159 183255500.0 6584.906
    14) humidity>=0.8322915 8 2719668.0 4539.750 *
    15) humidity< 0.8322915 151 145301700.0 6693.258
      30) month=2,3,4,5,7,11,12 76 65508410.0 6303.447 *
      31) month=6,8,9,10 75 56542530.0 7088.267 *
```

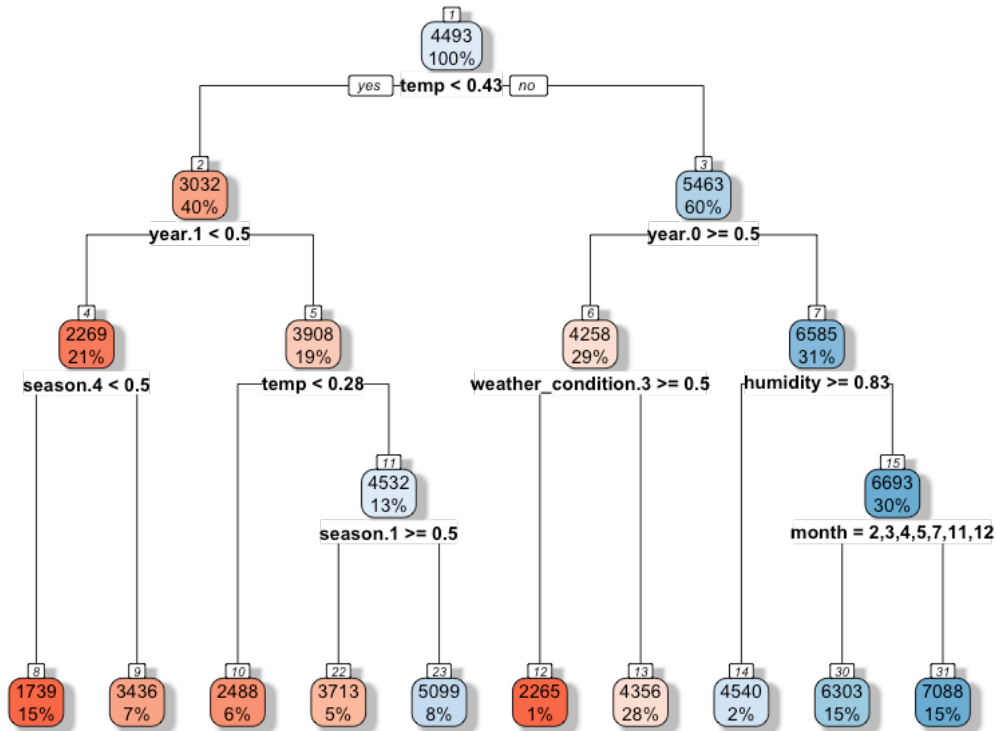
```
[276]: plot(original_fit_bshare$cptable[,4], main = "Cp for model selection", ylab = "cv error")

print(paste('Optimal tree size is:',min_cp))
```

```
[1] "Optimal tree size is: 10"
```



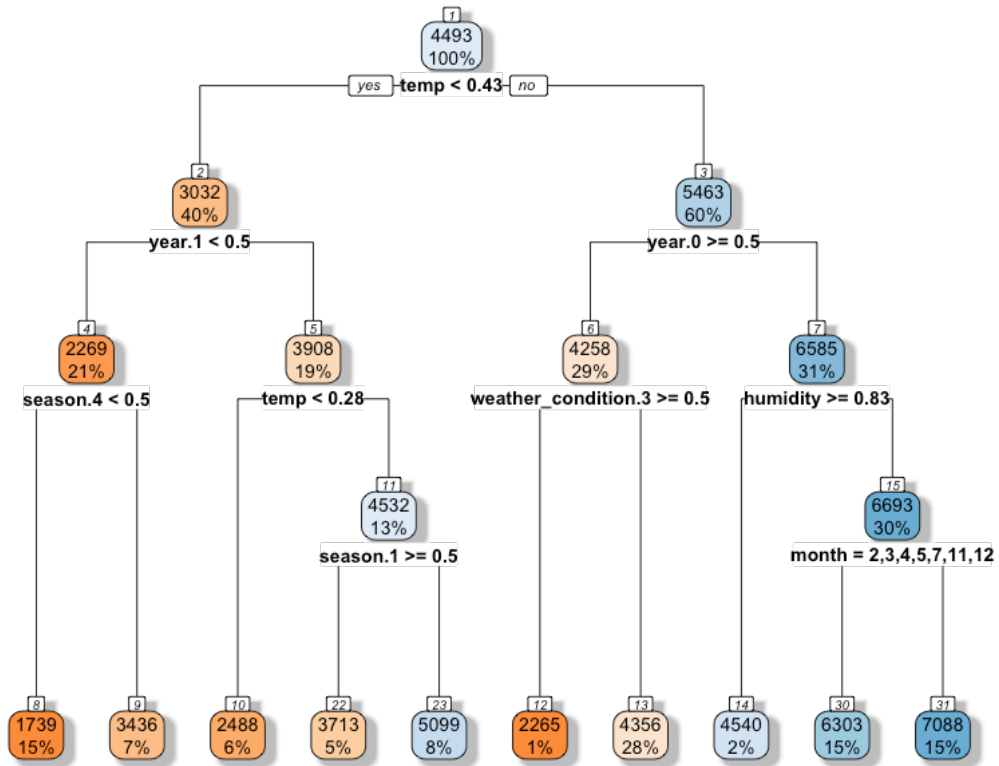
```
[277]: rpart.plot(pruned_fit_bshare, box.palette="RdBu", shadow.col="gray",  
↪ nn=TRUE, roundint=FALSE)
```



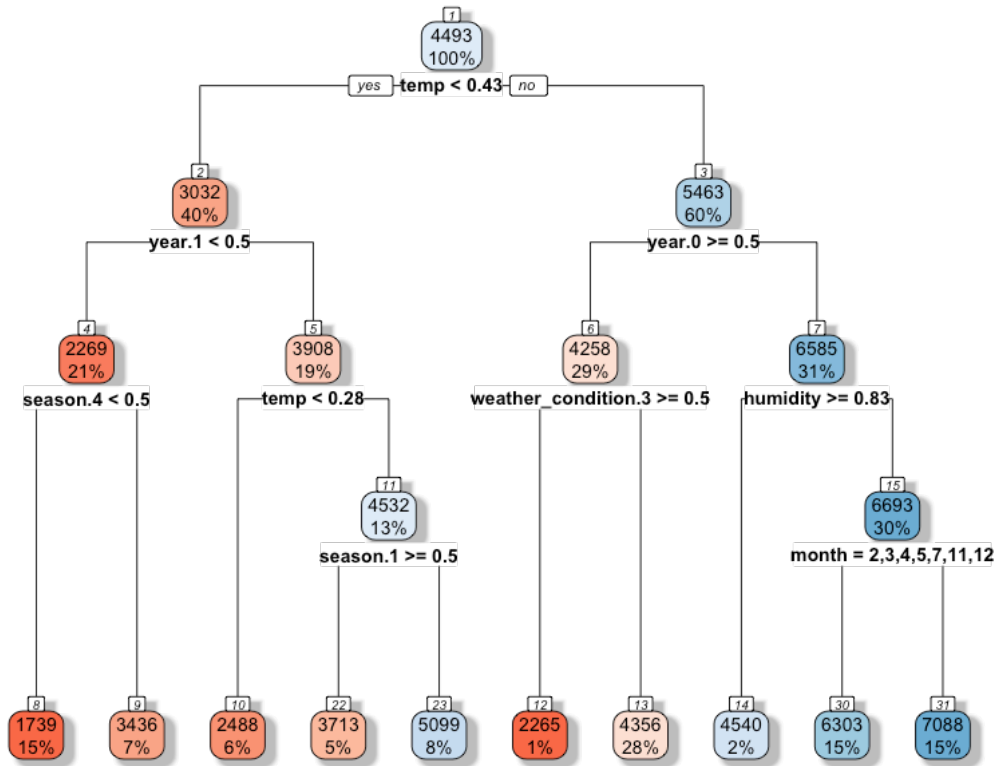
```
[244]: rpart.plot(original_fit_bshare, box.palette="OrBu", shadow.col="gray",
  ↪nn=TRUE,roundint=FALSE, main='Full Tree')

rpart.plot(pruned_fit_bshare, box.palette="RdBu", shadow.col="gray",
  ↪nn=TRUE,roundint=FALSE,main='Pruned Tree')
```

Full Tree



Pruned Tree



As you can see, pruning had no impact on the original tree. I believe this is because the tree was already in extremely good condition to begin with. Pruning is the process of removing leaves that do not considerably improve accuracy, i.e., to avoid overfitting. Since we deleted a few columns that seemed insignificant during the feature selection and preprocessing, the data is optimally balanced.

We can see that almost all terminal node have less than 10% of distribution of test samples. Only one node with decision 'no' has the distribution of 28%.

- 2.0.4** *(b) The board wants to identify good times to do “repair” and “tune-ups”. However, they want to minimize disruption of usage. You are asked to describe the “low yield” settings when the bikes are not being utilized (counts are low). Use your tree to advise on this matter.*

From the above pruned tree it is clear that the bikers count are low at the left most node of the tree. A warm summer day with minimal humidity would be ideal for the highest bike rental rates. We can infer that the quantity of bicycle rentals every day depends on a variety of parameters, including seasonal and weather-related ones. On the other hand, the fall (season 4) season, which runs from September through December, has the lowest bike rental rates, indicating that the two most significant factors influencing demand for bike sharing rentals are the season and the temperature. So the low yields can be maximized during these seasons.

3 Question 3

- 3.0.1** The Wage data (ISLR2) contains information related to demographics and earnings of males in the Mid-Atlantic area. Discretize “Wage” into three categories (low, med, high). You are asked to construct a classification tree to predict “Wage”.

- 3.0.2** *a) Divide the data into test and training and perform model selection to determine the optimal tree size. Comment on the performance of the tree. How many test samples are assigned to each terminal region?*

loading libraries

```
[180]: install.packages('ISLR')
install.packages('ggplot2')
install.packages('caret')
install.packages('tidyverse')
library(tidyverse)
library(ISLR)
library(ggplot2)
library(caret)
```

The downloaded binary packages are in
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

The downloaded binary packages are in
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

The downloaded binary packages are in
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

The downloaded binary packages are in
/var/folders/ch/lqq67w6x5px6fcc9cnbp457m0000gn/T//Rtmp2yueMP/downloaded_packages

Attaching packages
1.3.2 tidyverse

```

tibble 3.1.8      dplyr 1.0.10
tidyr  1.2.1      stringr 1.4.1
readr  2.1.3      forcats 0.5.2
purrr  0.3.4

Conflicts
tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag()     masks stats::lag()
purrr::lift()    masks caret::lift()
dplyr::select() masks MASS::select()

```

```
[888]: data(Wage)
       head(Wage,5)
```

		year <int>	age <int>	maritl <fct>	race <fct>	education <fct>	region <fct>
	231655	2006	18	1. Never Married	1. White	1. < HS Grad	2. Middle Atl
A data.frame: 5 × 11	86582	2004	24	1. Never Married	1. White	4. College Grad	2. Middle Atl
	161300	2003	45	2. Married	1. White	3. Some College	2. Middle Atl
	155159	2003	43	2. Married	3. Asian	4. College Grad	2. Middle Atl
	11443	2005	50	4. Divorced	1. White	2. HS Grad	2. Middle Atl

Now that we have seen that the variable “region” only has the value “Middle Atlantic,” we will remove it as it is unable to distinguish between any observations. Additionally, the predictor “logwage” will be removed because it has no significance with the dependent variable. We shall also eliminate that variable as we will discretize the “wage.” We can combine “3. Widowed”, “4. Divorced”, and “5. Separated” into a more populous level known as “3. Other” to increase the robustness of the models to be built because they all have somewhat similar meanings.

```
[889]: Wage = Wage %>%
       mutate(wage_category = case_when(wage < quantile(wage, 0.25) ~ "Low",
                                         wage >= quantile(wage, 0.25) & wage <=
                                         quantile(wage, 0.75) ~ "Medium",
                                         wage > quantile(wage, 0.75) ~ "High"))
       head(Wage,5)
```

		year <int>	age <int>	maritl <fct>	race <fct>	education <fct>	region <fct>
	231655	2006	18	1. Never Married	1. White	1. < HS Grad	2. Middle Atl
A data.frame: 5 × 12	86582	2004	24	1. Never Married	1. White	4. College Grad	2. Middle Atl
	161300	2003	45	2. Married	1. White	3. Some College	2. Middle Atl
	155159	2003	43	2. Married	3. Asian	4. College Grad	2. Middle Atl
	11443	2005	50	4. Divorced	1. White	2. HS Grad	2. Middle Atl

```
[890]: Wage$region <- NULL
       Wage$wage <- NULL
       Wage$logwage <- NULL
       levels(Wage$maritl)[3:5] <- "3. Other"
```

Checking for missing values

```
[891]: any(is.na(Wage))
```

FALSE

Split the dataset into train and test

```
[892]: set.seed(2021)
partition <- createDataPartition(y = as.factor(Wage$wage_category), p = .7, list_
  ↪ = FALSE)
data.train <- Wage[partition, ]
data.test <- Wage[-partition, ]
head(data.train, 5)
```

	year	age	maritl	race	education	jobclass	
	<int>	<int>	<fct>	<fct>	<fct>	<fct>	
A data.frame: 5 × 9	86582	2004	24	1. Never Married	1. White	4. College Grad	2. Information
	161300	2003	45	2. Married	1. White	3. Some College	1. Industrial
	11443	2005	50	3. Other	1. White	2. HS Grad	2. Information
	376662	2008	54	2. Married	1. White	4. College Grad	2. Information
	450601	2009	44	2. Married	4. Other	3. Some College	1. Industrial

```
[893]: data.train$wage_category <- factor(data.train$wage_category)
```

```
[894]: head(data.train, 5)
```

	year	age	maritl	race	education	jobclass	
	<int>	<int>	<fct>	<fct>	<fct>	<fct>	
A data.frame: 5 × 9	86582	2004	24	1. Never Married	1. White	4. College Grad	2. Information
	161300	2003	45	2. Married	1. White	3. Some College	1. Industrial
	11443	2005	50	3. Other	1. White	2. HS Grad	2. Information
	376662	2008	54	2. Married	1. White	4. College Grad	2. Information
	450601	2009	44	2. Married	4. Other	3. Some College	1. Industrial

```
[895]: original_fit_wage = rpart(wage_category ~. , data = data.train, method =
  ↪ "class", control = rpart.control(minbucket=5, cp=0.0005, maxdepth=7), parms_
  ↪ = list(split = "gini"))

original_fit_wage
```

n= 2102

node), split, n, loss, yval, (yprob)

* denotes terminal node

1) root 2102 1038 Medium (0.24881066 0.24500476 0.50618459)

2) education=4. College Grad,5. Advanced Degree 795 427 High (0.46289308 0.
 ↪10440252 0.43270440)

4) age>=33.5 632 300 High (0.52531646 0.09177215 0.38291139)
 8) education=5. Advanced Degree 259 88 High (0.66023166 0.06177606 0.
 ↪27799228)
 16) health_ins=1. Yes 209 61 High (0.70813397 0.02870813 0.26315789)
 32) maritl=2. Married,3. Other 185 46 High (0.75135135 0.02702703 0.
 ↪22162162)
 64) health=2. >=Very Good 149 30 High (0.79865772 0.02013423 0.
 ↪18120805) *
 65) health=1. <=Good 36 16 High (0.55555556 0.05555556 0.38888889)
 130) year>=2006.5 15 4 High (0.73333333 0.06666667 0.20000000) *
 131) year< 2006.5 21 10 Medium (0.42857143 0.04761905 0.
 ↪52380952) *
 33) maritl=1. Never Married 24 10 Medium (0.37500000 0.04166667 0.
 ↪58333333)
 66) jobclass=2. Information 19 9 Medium (0.47368421 0.00000000 0.
 ↪52631579)
 132) age< 49 13 5 High (0.61538462 0.00000000 0.38461538) *
 133) age>=49 6 1 Medium (0.16666667 0.00000000 0.83333333) *
 67) jobclass=1. Industrial 5 1 Medium (0.00000000 0.20000000 0.
 ↪80000000) *
 17) health_ins=2. No 50 27 High (0.46000000 0.20000000 0.34000000)
 34) age>=37.5 44 23 High (0.47727273 0.22727273 0.29545455)
 68) age< 62 38 19 High (0.50000000 0.26315789 0.23684211)
 136) race=2. Black,3. Asian 6 2 Low (0.33333333 0.66666667 0.
 ↪00000000) *
 137) race=1. White 32 15 High (0.53125000 0.18750000 0.28125000)↪
 ↪*
 69) age>=62 6 2 Medium (0.33333333 0.00000000 0.66666667) *
 35) age< 37.5 6 2 Medium (0.33333333 0.00000000 0.66666667) *
 9) education=4. College Grad 373 203 Medium (0.43163539 0.11260054 0.
 ↪45576408)
 18) health_ins=1. Yes 295 153 High (0.48135593 0.07118644 0.44745763)
 36) age< 64.5 289 147 High (0.49134948 0.06920415 0.43944637)
 72) maritl=2. Married 227 108 High (0.52422907 0.06167401 0.
 ↪41409692)
 144) year< 2008.5 194 88 High (0.54639175 0.06185567 0.39175258)↪
 ↪*
 145) year>=2008.5 33 15 Medium (0.39393939 0.06060606 0.
 ↪54545455) *
 73) maritl=1. Never Married,3. Other 62 29 Medium (0.37096774 0.
 ↪09677419 0.53225806) *
 37) age>=64.5 6 1 Medium (0.00000000 0.16666667 0.83333333) *
 19) health_ins=2. No 78 40 Medium (0.24358974 0.26923077 0.48717949)
 38) year< 2005.5 38 22 Low (0.23684211 0.42105263 0.34210526)
 76) age< 39.5 9 4 High (0.55555556 0.33333333 0.11111111) *
 77) age>=39.5 29 16 Low (0.13793103 0.44827586 0.41379310)

```

154) health=1. <=Good 11    4 Low (0.09090909 0.63636364 0.
↳27272727) *
155) health=2. >=Very Good 18    9 Medium (0.16666667 0.33333333 0.
↳50000000) *
39) year>=2005.5 40    15 Medium (0.25000000 0.12500000 0.62500000)
78) age>=55.5 7    4 High (0.42857143 0.28571429 0.28571429) *
79) age< 55.5 33    10 Medium (0.21212121 0.09090909 0.69696970) *
5) age< 33.5 163    61 Medium (0.22085890 0.15337423 0.62576687)
10) race=3. Asian 23    11 High (0.52173913 0.08695652 0.39130435) *
11) race=1. White,2. Black,4. Other 140    47 Medium (0.17142857 0.
↳16428571 0.66428571)
22) age< 24.5 15    7 Low (0.06666667 0.53333333 0.40000000)
44) year< 2004.5 7    2 Low (0.00000000 0.71428571 0.28571429) *
45) year>=2004.5 8    4 Medium (0.12500000 0.37500000 0.50000000) *
23) age>=24.5 125    38 Medium (0.18400000 0.12000000 0.69600000)
46) year>=2004.5 87    31 Medium (0.25287356 0.10344828 0.64367816)
92) age>=32.5 15    8 High (0.46666667 0.06666667 0.46666667)
184) maritl=2. Married 8    3 High (0.62500000 0.00000000 0.
↳37500000) *
185) maritl=1. Never Married 7    3 Medium (0.28571429 0.14285714
↳0.57142857) *
93) age< 32.5 72    23 Medium (0.20833333 0.11111111 0.68055556) *
47) year< 2004.5 38    7 Medium (0.02631579 0.15789474 0.81578947) *
3) education=1. < HS Grad,2. HS Grad,3. Some College 1307    587 Medium (0.
↳11859220 0.33052793 0.55087988)
6) age< 25.5 134    26 Low (0.01492537 0.80597015 0.17910448)
12) maritl=1. Never Married,3. Other 115    14 Low (0.00000000 0.87826087
↳0.12173913) *
13) maritl=2. Married 19    9 Medium (0.10526316 0.36842105 0.52631579)
26) education=1. < HS Grad,2. HS Grad 13    7 Low (0.15384615 0.
↳46153846 0.38461538) *
27) education=3. Some College 6    1 Medium (0.00000000 0.16666667 0.
↳83333333) *
7) age>=25.5 1173    477 Medium (0.13043478 0.27621483 0.59335038)
14) health_ins=2. No 400    211 Medium (0.06000000 0.46750000 0.47250000)
28) education=1. < HS Grad 81    23 Low (0.00000000 0.71604938 0.
↳28395062)
56) year< 2005.5 36    6 Low (0.00000000 0.83333333 0.16666667)
112) race=1. White,2. Black 31    3 Low (0.00000000 0.90322581 0.
↳09677419) *
113) race=3. Asian,4. Other 5    2 Medium (0.00000000 0.40000000 0.
↳60000000) *
57) year>=2005.5 45    17 Low (0.00000000 0.62222222 0.37777778)
114) age< 35 11    1 Low (0.00000000 0.90909091 0.09090909) *
115) age>=35 34    16 Low (0.00000000 0.52941176 0.47058824)
230) race=2. Black,3. Asian,4. Other 12    2 Low (0.00000000 0.
↳83333333 0.16666667) *

```

231) race=1. White 22 8 Medium (0.00000000 0.36363636 0.
↳63636364) *

29) education=2. HS Grad,3. Some College 319 153 Medium (0.07523511 0.
↳40438871 0.52037618)

58) race=2. Black,3. Asian,4. Other 63 28 Low (0.07936508 0.
↳55555556 0.36507937)

116) year>=2003.5 51 20 Low (0.05882353 0.60784314 0.33333333) *
117) year< 2003.5 12 6 Medium (0.16666667 0.33333333 0.50000000) *

59) race=1. White 256 113 Medium (0.07421875 0.36718750 0.55859375)
118) health=1. <=Good 93 48 Medium (0.04301075 0.47311828 0.
↳48387097)

236) year< 2003.5 9 1 Low (0.11111111 0.88888889 0.00000000) *
237) year>=2003.5 84 39 Medium (0.03571429 0.42857143 0.
↳53571429) *

119) health=2. >=Very Good 163 65 Medium (0.09202454 0.30674847 0.
↳60122699) *

15) health_ins=1. Yes 773 266 Medium (0.16688228 0.17723157 0.65588616)
30) education=3. Some College 288 92 Medium (0.24305556 0.07638889 0.
↳68055556)

60) maritl=2. Married 205 72 Medium (0.29756098 0.05365854 0.
↳64878049) *

61) maritl=1. Never Married,3. Other 83 20 Medium (0.10843373 0.
↳13253012 0.75903614)

122) year< 2004.5 24 11 Medium (0.16666667 0.29166667 0.54166667)
244) race=2. Black,4. Other 6 3 Low (0.33333333 0.50000000 0.
↳16666667) *

245) race=1. White,3. Asian 18 6 Medium (0.11111111 0.22222222
↳0.66666667) *

123) year>=2004.5 59 9 Medium (0.08474576 0.06779661 0.84745763) *
31) education=1. < HS Grad,2. HS Grad 485 174 Medium (0.12164948 0.
↳23711340 0.64123711)

62) maritl=2. Married,3. Other 419 145 Medium (0.14081146 0.20525060
↳0.65393795) *

63) maritl=1. Never Married 66 29 Medium (0.00000000 0.43939394 0.
↳56060606)

126) jobclass=1. Industrial 41 20 Low (0.00000000 0.51219512 0.
↳48780488)

252) health=2. >=Very Good 26 10 Low (0.00000000 0.61538462 0.
↳38461538) *

253) health=1. <=Good 15 5 Medium (0.00000000 0.33333333 0.
↳66666667) *

127) jobclass=2. Information 25 8 Medium (0.00000000 0.32000000 0.
↳68000000)

254) age>=42 13 6 Low (0.00000000 0.53846154 0.46153846) *
255) age< 42 12 1 Medium (0.00000000 0.08333333 0.91666667) *


```
[896]: prop.table(table(data.train$wage_category))
```

```
      High      Low      Medium
0.2488107 0.2450048 0.5061846
```


```
[897]: original_fit_wage$cptable
plotcp(original_fit_wage)
```

A matrix: 14 × 5 of type dbl

	CP	nsplit	rel error	xerror	xstd
1	0.0558766859	0	1.0000000	1.0000000	0.02208288
2	0.0168593449	3	0.8323699	0.8545279	0.02181405
3	0.0115606936	5	0.7986513	0.8265896	0.02170902
4	0.0091522158	6	0.7870906	0.8256262	0.02170508
5	0.0072254335	8	0.7687861	0.8265896	0.02170902
6	0.0048169557	10	0.7543353	0.7976879	0.02158173
7	0.0038535645	11	0.7495183	0.7957611	0.02157256
8	0.0028901734	13	0.7418112	0.8082852	0.02163062
9	0.0024084778	16	0.7331407	0.8111753	0.02164351
10	0.0019267823	20	0.7235067	0.8121387	0.02164776
11	0.0014450867	28	0.7080925	0.8246628	0.02170112
12	0.0009633911	34	0.6994220	0.8371869	0.02175092
13	0.0006422608	41	0.6926782	0.8487476	0.02179376
14	0.0005000000	47	0.6888247	0.8506744	0.02180060




```
[898]: rpart.plot(original_fit_wage, box.palette="OrBu", shadow.col="gray",
↪ nn=TRUE, roundint=FALSE, main='Full Tree')
```

```
[899]: plot(original_fit_wage$cptable[,4], main = "Cp for model selection", ylab = "cv_
↪error")

print(paste('Optimal tree size is:',which.min(original_fit_wage$cptable[,4])))
```

```
[1] "Optimal tree size is: 7"
```



```
[902]: min_cp = original_fit_wage$cptable[which.  

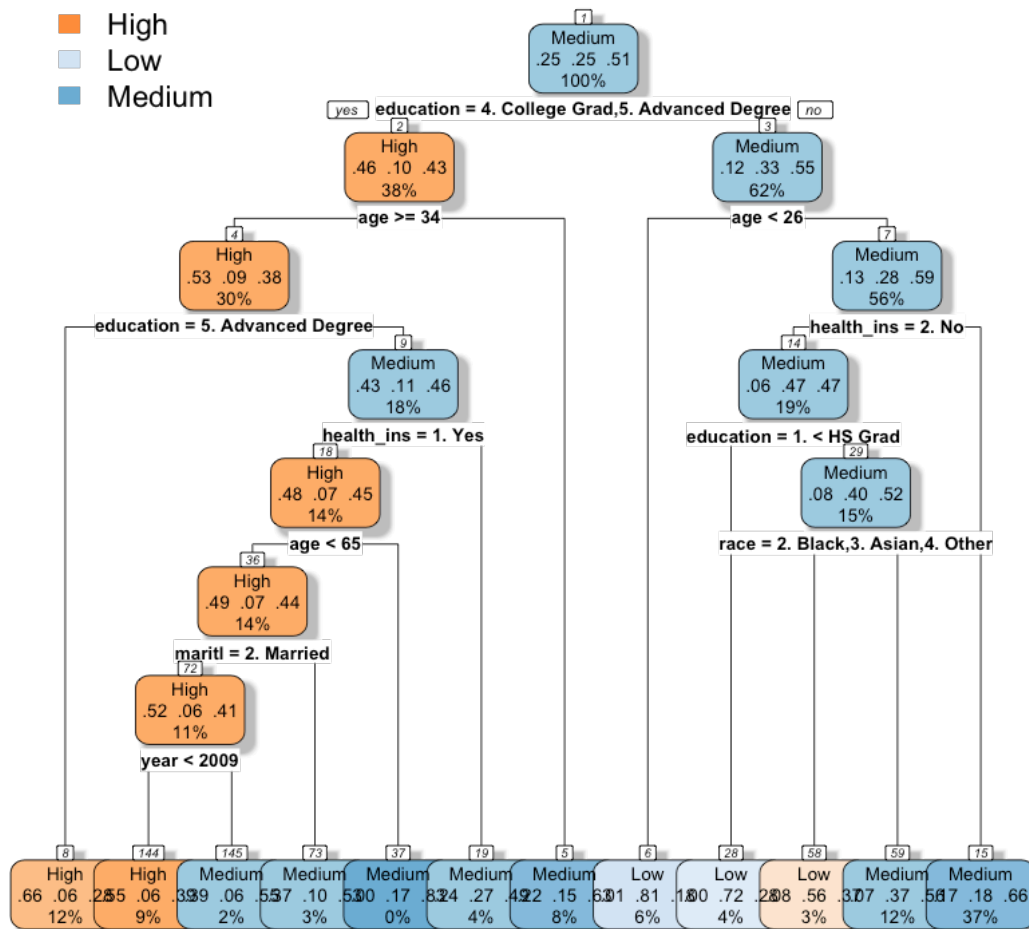
  ↪min(original_fit_wage$cptable[, "xerror"]), "CP"]  

pruned_fit_wage <- prune(original_fit_wage, cp = min_cp)  

rpart.plot(pruned_fit_wage, tweak=1.5, box.palette="OrBu", shadow.col="gray",  

  ↪nn=TRUE, roundint=FALSE, main='Pruned Tree')
```

Pruned Tree



The tree shows that the most important predictor of whether a worker will make a high income is education. The factors that follow schooling in terms of relevance are health indicators, age, and marital status. The pruned tree has 12 terminal nodes and 11 splits. Although it appears simpler than the original tree, the pruned tree is nonetheless complex.

3.0.3 *b) Model interpretation – you are asked to characterize the high wage earners. What can you say about this group (based on the tree)?*

The pruned tree has 12 terminal nodes and 11 splits. Tree is considerably easier to understand and comprehend than the original tree. The first division is made based on education, which can be thought of as the main predictor. The following divisions include health insurance, marital status, age, year, and education once again. The first division separates areas with high education levels to the left and relatively low education levels to the right. Low/medium earners are the anticipated class for those with low levels of education. High earners are the anticipated class for those with

high education levels. The race of Black, Asian, and others is used as the cutoff level for further partitioning of Node 9. Workers of these races are categorized as low earners. Similar to this, Node 8 is further partitioned using a cutoff level of 65 years old. Workers that are younger than this age threshold are categorized as high earners.

[]: